

radar camera fusion

我们的方法称为CenterFusion，它首先使用中心点检测网络通过在图像上识别对象的中心点来检测对象。

然后它使用一种新颖的基于视锥的方法来解决关键数据关联问题，从而将雷达检测结果与相对应对象的中心点相关联。

相关的radar检测用于生成基于radar的特征图以补充图像特征，并回归到诸如深度，旋转和速度之类的对象属性。

radar使用多普勒效应快速准确的确定物体的速度。

视锥体关联机制

使用对象的2d边框估计其深度和大小，为该对象创建3d兴趣区域，视锥体。

在roi内离图像中心点最近的点，关联。

$$Y_{qc} = \max_i \exp\left(-\frac{(p_i - q)^2}{2\delta_i^2}\right)$$

其中 δ_i 是尺寸自适应标准偏差。

For autonomous robots to navigate a complex environment , it is crucial 至关重要的 to understand the surrounding scene both geometrically and semantically.Modern autonomous robots employ multiple sets of sensors, including lidars, radars, and cameras. Managing the different reference frames and characteristics of the sensors, and merging their observations into a single representation complicates perception.

Our approach consists of an ensemble of neural networks which take in sensor data from different modalities 形式 and transform them into a single common top-down semantic grid representation.We find representation favourable as it is agnostic无关 to sensor-specific reference frames and captures both the semantic and geometric information for the surrounding scene.

Because the modalities share a single output representation,they can be easily aggregated 汇总 to produce a fused output.In this work we predict short-term semantic grids but the framework can be extended to other tasks.This approach offers a simple, extensible , end to end approach for multi-modal perception and prediction.

The benefits of a shared top-down representation across modalities are threefold. First, it is an interpretable representation that better facilitates促进 debugging调试 and reasoning推理 about inherent固有 failure modes故障模式 of each modality.Second it is independent of any particular sensors characteristics and so is easily extensible for adding new modalities.Finally , it simplifies the task of late fusion by sharing a spatial空间的 representation in a succinct manner.

In this work we present a novel end-to-end framework that predicts the top-down view of the current scene(\$t_0\$) as well as multiple timesteps into the future.The pipeline consists of a convolutional neural network for each of three sensor modalities : lidar, radar , camera.Each sensor modality predicts a sequence

of top-down semantic grids, then these outputs are fused to produce a single output grid. We explore fusing using two different aggregation mechanisms.

Estimating 3D orientation and translation of objects is essential for infrastructure-less .In case of monocular vision, successful methods have been mainly based on two ingredients因素: (1) a network generating 2D region proposals(2D区域提案)。 (2) a R-CNN structure predicting 3D object pose by utilizing the acquired regions of interest. We argue that the 2D detection network is redundant多余的 and introduces non-negligible 不可忽略的 noise for 3D detection. Hence, we propose a novel 3D object detection method, named SMOKE, in this paper that predicts a 3D bounding box for each detected object by combining a single keypoint estimate with regressed 3D variables. As a second contribution, we propose a multi-step disentangling approach for constructing the 3D bounding box ,which significantly improves both training convergence收敛 and detection accuracy. In constrast to previous 3D detection techniques, our method does not require complicated pre/post-processing.extra data, and a refinement细化 stage. Despite of its structural结构 simplicity, our proposed SMOKE network outperforms all existing monocular 3D detection methods on the KITTI dataset , giving the best state-of-the-art result on both 3D objection dection and bird's eye view evaluation . The code will be made publicly available.

Vison based object detection is an essential ingredient成分 of autonomous vehicle perception 洞察力 of autonomous vehicle perception and infrastructure基础设施 less robot navigation in general. This type of detection methods are used to perceive 感知 the surrounding environment by detecting and classifying object instances实例 into categories 类别 and identifying their locations and orientations. Recent developments in 2D object detection have achieved promising performance有前途的 on both detection accuracy and speed. In constrast, 3D object detection have proven to be a more challenging task as it aims to estimate pose and location for each object simulataneously.

Currently, the most successful 3D object detection methods heavily depend on Lidar point cloud.




Figure 1. SMOKE directly predicts the 3D projected keypoint and 3D regression parameters on a single image. The whole network is trained end-to-end in a single stage.

or LIDAR-Image fusion information.(features learned from the point cloud are key components of the detection network).However , LIDAR sensors are extremely expensive , have a short service life time and too heavy for autonomous robots. Hence Lidars are currently not considered to be economical流行的 to support autonomous vehicle operations . Alternatively , cameras are cost-effective, easily mountable and light-weight solutions for 3D object detection with long expected service time. Unlike lidar sensors , a single camera in itself cannot obtain sufficient spatial information for the whole environment as single RGB images can not supply object location information or dimensional contour 轮廓 in the real world. While

binocular 双目 vision restores the missing spatial information. in many robot applications, especially UAVs , it is difficult to realize biocular vision. Hence , it is desirable to perform 3D detection on a monocular image even if it is a more difficult and challenging task.

To enhance performance, geometry reasoning 几何推理 synthetic data 综合数据 and post 3D-2D processing have also been used to improve 3D object detection on single image. By the knowledge of the authors , no reliable monocular 3D detection method has been introduced so far to learn 3D information directly from image plane avoiding the performance decrease that is inevitable 不可避免的 with multi-stage method.

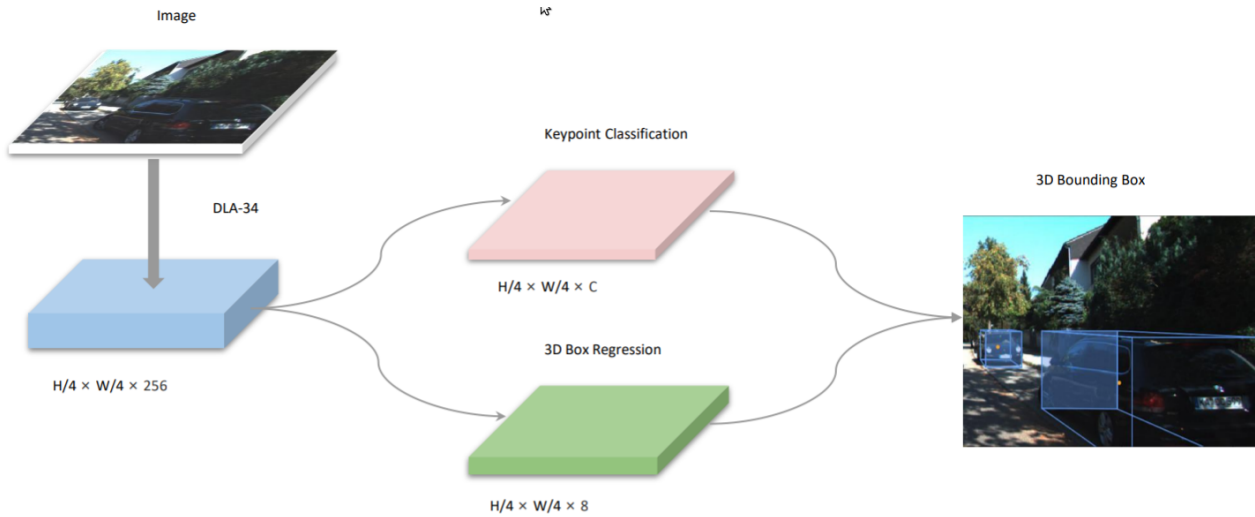


Figure 2. **Network Structure of SMOKE.** We leverage DLA-34 [41] to extract features from images. The size of the feature map is 1:4 due to downsampling by 4 of the original image. Two separate branches are attached to the feature map to perform keypoint classification (pink) and 3D box regression (green) jointly. The 3D bounding box is obtained by combining information from two branches.

In this paper we propose an innovative 创新 single-stage 3D object detection method that pairs each object with a single keypoint. We argue and later show that a 2D detection, which introduces nonnegligible noise in 3D parameter estimation, is redundant 多余的 to perform 3D object detection. Furthermore 2D information can be naturally obtained if 3D variables and camera intrinsic matrix are already known. Consequently, our designed network eliminate 排除 the 2D detection branch and estimates the projected 3D points on the image plane instead. A 3d parameter regression branch is added in parallel. This design results in a simple network structure with two estimation threads. Rather than regressing variables in a separate method by use multiple loss functions, we transform these variables together with projected keypoint to 8 corner representation of 3D boxes and regress them with a unified 统一的 loss function. As in most single state 2D object detection algorithms, our 3D detection approach only contains one classification and regression branch. Benefiting from the simple structure , the network exhibits 展示 improved accuracy in learning 3D variables, has better convergence and less overall computational needs.

Second contribution of our work is a multi-step disentanglement 纠缠 approach for 3D bounding box regression. Since all the geometry information is grouped into the parameter, it is difficult for the network to learn each variable accurately in a unified way. Our proposed method isolates 分离 the contribution of each parameter in both the 3D bounding box encoding phase and regression loss function , which significantly helps to train the whole network effectively.

Our contribution is summarized as follows:

- We propose a one-stage monocular 3D object detection with a simple architecture that can precisely learn 3D geometry in an end-to-end fashion.

- We provide a multistep distanglement approach to improve the convergence收敛 of 3D parameters and detection accuracy.
- The result method outperforms all existing state-of-art monocular 3D object detection algorithms on the chanllenging KITTI dataset at the submission date November 2019.



Figure 3. Visualization of difference between 2D center points (red) and 3D projected points (orange). Best viewed in color.

We formulate定义 the monocular 3D object detection problem as follow:

given a single RGB image $I \in \mathbb{R}^{W \times H \times 3}$, with W the width and H the height of the image, find for each present object its category类别 label C and its 3D bounding box B , where the latter is parameterized by 7 variables $(h, \omega, l, x, y, z, \theta)$. Here, (h, ω, l) represent the height, width and length of each object in meters, and (x, y, z) is the coordinates (in meters) of the object center in the camera coordinate frame. Variable θ is the yaw orientation of the corresponding cubic box. The roll and pitch angles are set to zero by following the Kitti annotation. Additionally, we take the mild温和 assumption that the camera intrinsic matrix K is known for both training and inference.

Smoke approach

In this section, we describe the smoke network that directly estimates 3D bounding boxes for detected object instances from monocular imagery. In constrast to previous techniques that leverage 杠杆作用 2D proposals to predict 3D bounding box, our method can detect 3D information with a simple single stage. The propose method can be divided into three parts: (1) backbone, (2) 3D detection (3) loss function. First, we briefly discuss the backbone for feature extraction, followed by the instruction of the the 3D detection network consisting of two separated branches. Finally, we discuss the loss function design and the multi-step disentanglement to compute the regression loss. The overview of the network structure is depicted in Fig.2

backbone

We use a hierarchical layer fusion network DLA-34 as the backbone to extract features since it can aggregate information across different layers. Following the same structure as in [], all the hierarchical aggregation connections are replaced by a Deformable可变形的 convolution network (DCN). The output feature map is downsample 4 times with respect to the original image. Compared with the original implementation, we replace all BatchNorm(BN) operation with GroupNorm(GN) since it has been proven to

be less sensitive to batch size and more robust to training noise. We also use this technique in the two prediction branches.




Figure 4. Relation of the observation angle α_x and α_z . α_x is provided in KITTI, while α_z is the value we choose to regress.

which will be discussed later. This adjustment not only improves detection accuracy, but it also reduces considerably the training time. Later, we provide performance comparison of BN and GN to demonstrate these properties.

3D detection network

Keypoint branch: We define the keypoint estimation network similar to such that each object is represented by one specific keypoint. Instead of identifying the center of a 2D bounding box, the key point is defined as the projected 3D center of the object on the image plane. The comparison between 2D center points and 3D projected points is visualized in Fig3. The projected keypoints allow to fully recover 3D location for each object with camera parameters. Let $[x, y, z]^T$ represent the 3D center of each object in the camera frame. The projection of 3D points of points $[x_c, y_c]^T$ on the image plane can be obtained with the camera intrinsic matrix K in a homogeneous form:

$$\begin{bmatrix} z \\ x_c \\ y_c \end{bmatrix} = K^{-1} \begin{bmatrix} 3 \\ x \\ y \\ z \end{bmatrix} \quad (1)$$

For each ground truth keypoint, its corresponding downsampled location on the feature map is computed and distributed using a Gaussian Kernel following. The standard deviation is allocated based on the 3D bounding boxes of the ground truth projected to the image plane. Each 3D box on the image is represented by 8 2D points $[x_{b,1\sim 8}, y_{b,1\sim 8}]^T$ and the standard deviation is computed by the smallest 2D box with $\{x_b^{min}, y_b^{min}, x_b^{max}, y_b^{max}\}$ that encircles the 3D box.

Regression Branch

Each key point on the heatmap. Similar to other monocular 3D detection framework. The 3D information is encoded as an 8-tuple 元组

$$\tau = \begin{bmatrix} \delta_z & \delta_{x_c} & \delta_{y_c} & \delta_h & \delta_w & \delta_l & \sin(\alpha) & \cos(\alpha) \end{bmatrix}^T$$

Here δ_z denotes the depth offset, $\delta_{x_c}, \delta_{y_c}$ is the discretization 离散化 offset due to downsampling. $\delta_h, \delta_w, \delta_l$ denotes the residual 剩余 dimension. $\sin(\alpha), \cos(\alpha)$ is the vectorial representation of the rotational angle α . We encode all variables to be learnt in residual 残差 representation to reduce the learning interval and ease the training task. The size of feature map for regression results in $S_r \in \mathbb{R}^{\frac{H}{R} \times \frac{W}{R} \times 8}$. Inspired by the lifting transformation 提升转型 described in [], we introduce a similar operation F that converts projected 3D points to a 3D bounding box $B = F(\tau) \in \mathbb{R}^{3 \times 8}$. For each object, its depth z can be recovered by pre-defined scale and shift parameters σ_z and μ_z as $z = \mu_z + \delta_z \sigma_z$.

Given the object depth z , the location for each object in the camera frame can be recovered by using its discretized 离散的 projected centroid $[x_c, y_c]^T$ on the image plane and the downsampling offset

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = K^{-1} \begin{bmatrix} 3 \\ x_c + \delta_{x_c} \\ y_c + \delta_{y_c} \\ z \end{bmatrix}$$

This equation is inverse of Eq.(1). In order to retrieve 取回 object dimensions $[h, w, l]^T$, we use a pre-calculated category-wise 按类别 average dimension $[\bar{h}, \bar{w}, \bar{l}]^T$ computed over the whole dataset. Each object dimension can be recovered by using the residual dimension offset $[\delta_h, \delta_w, \delta_l]^T$:

$$\begin{bmatrix} h \\ w \\ l \end{bmatrix} = \begin{bmatrix} \overline{h} \\ \overline{w} \\ \overline{l} \end{bmatrix} \begin{bmatrix} e^{\delta_h} \\ e^{\delta_w} \\ e^{\delta_l} \end{bmatrix} \quad (4)$$

Inspired by [], we choose to regress the observation angle α instead of the yaw rotation θ for each object. We further change the observation angle with respect to the object head α_x , instead of the commonly

used observation angle value α_z , by simply adding $\frac{\pi}{2}$. The difference between these two angles is shown in Figure 4. Moreover, each α is encoded as the vector $[\sin(\alpha), \cos(\alpha)]^T$. The yaw angle θ can be obtained by utilizing α_z and the object location: $\theta = \alpha_z + \arctan\left(\frac{x}{z}\right)$ (5). Finally, we can construct the 8 corners of the 3D bounding box in the camera frame by using the yaw rotation matrix R_θ , object dimensions $[h, w, l]^T$ and location

$$[x \quad y \quad z]^T: B = R_{\theta} \begin{bmatrix} h/2 \\ w/2 \\ l/x \\ y/z \end{bmatrix} \quad (6)$$

Loss function

We employ the penalty-reduce focal loss 惩罚减少焦点损失 [1] in a point-wise manner on the downsampled heatmap. Let $s_{i,j}$ be the predicted score at the heatmap location (i,j) and $y_{i,j}$ be the ground-truth value of each point assigned by Gaussian Kernel. Define $\breve{y}_{i,j}$ and $\breve{s}_{i,j}$ as : $\breve{y}_{i,j} = \begin{cases} 0, & \text{if } y_{i,j} = 1 \\ y_{i,j}, & \text{otherwise} \end{cases}$. $\breve{s}_{i,j} = \begin{cases} s_{i,j}, & \text{if } y_{i,j} = 1 \\ 1 - s_{i,j}, & \text{otherwise} \end{cases}$

For simplicity, we only consider a single object class here. Then, the classification loss function is constructed as $L_{cls} = -\frac{1}{N} \sum_{h,w} (1 - \breve{y}_{i,j})^{\beta} (1 - \breve{s}_{i,j})^{\alpha} \log(\breve{s}_{i,j})$ where α and β are tunable hyper-parameters and N is the number of keypoints per image. The term $(1 - y_{i,j})$ corresponds to penalty reduction for points around the groundtruth location.

Regression Loss

We regard the 8D tuple τ to construct the 3D bounding box for each object. We also add channel-wise action to the regressed parameters of dimension and orientation at each feature map location to preserve consistency 一致性. The activation functions for the dimension and the orientation are chosen to be the sigmoid function σ and the l_2 norm, respectively: $\begin{bmatrix} \Delta h \\ \Delta w \\ \Delta l \end{bmatrix} = \sigma \left(\begin{bmatrix} O_h \\ O_w \\ O_l \end{bmatrix} \right) - \frac{1}{2}$, $\begin{bmatrix} \sin(a) \\ \cos(a) \end{bmatrix} = \begin{bmatrix} O_{\sin} \\ O_{\cos} \end{bmatrix} / \sqrt{O_{\sin}^2 + O_{\cos}^2}$. Here O stands for the specific output of network. By adopting the keypoint lifting transformation introduced in Sec.4.2, we define the 3D bounding box regression loss as the l_1 distance between the predicted transform \hat{B} and the ground truth B : $L_{reg} = \frac{\lambda}{N} ||\hat{B} - B||_1$ where λ is the scaling factor. This is used to ensure that neither the classification, nor the regression dominates 主导 the other. The disentangling transformation of loss has been proven to be an effective dynamic method to optimize 3D regression loss functions in [31]. Following this design, we extend the concept of loss disentanglement into a multi-step form. In Eq(3), we use the projected 3D groundtruth points on the image plane $[x_c \quad y_c]^T$ with the network predicted discretization offset $[\hat{\Delta x_c} \quad \hat{\Delta y_c}]^T$ and depth \hat{z} to retrieve the location $[\hat{x} \quad \hat{y} \quad \hat{z}]^T$ of each object. In Eq(5), we use the groundtruth location $[x \quad y \quad z]^T$ and the predicted observation angle \hat{a} to construct the estimated yaw orientation $\hat{\theta}$. The 8 corners representation of the 3D bounding box is also isolated 隔离 into three different groups following the concept of disentanglement, namely orientation, dimension and location. The final loss function can be represented by: $L = L_{cls} + \sum_{i=1}^3 L_{reg}(\hat{B}_i)$ where i represents the number of groups we define in the 3D regression branch. The multi-step disentangling transformation divides the contribution of each parameter group to the final loss. In Sec.5.2, we show that this method significantly improves detection accuracy.

Implementation

In this section, we discuss the implementation of our proposed methodology in detail together with selection of the hyperparameters.

Preprocessing : we avoid applying any complicated preprocessing method on the dataset. Instead, we only eliminate 排除 objects whose 3D projected center point on the image plane is out of the image range. Note that the total number of projected center points outside the image boundary for the car instance is 1582. This accounts for only the 5.5% of the entire set of 28742 labeled cars.

Data Augmentation : Data augmentation techniques we used are random horizontal flip 水平翻转, random scale and shift. The scale ratio is set to 9 steps from 0.6 to 1.4, and the shift ratio is set to 5 steps from -0.2 to 0.2. Note that the scale and shift augmentation methods are only used for heatmap classification since the 3D information becomes inconsistent 不一致 with data augmentation.

Hyperparameter Choice : In the backbone, the group number for GroupNorm is set to 32. For channels less than 32, it is set to be 16. For Eq.7, we set $\alpha=2$ and $\beta=4$ in all experiments. Based on [31], the reference car size and depth statistic 统计数据 we use are $[\overline{h} \quad \overline{w} \quad \overline{l}]^T = [1.63 \quad 1.53 \quad 3.88]^T$ and $[\mu_z \quad \sigma_z]^T = [28.01 \quad 16.32]^T$ (measured in meters).

Training : Our optimization schedule 优化表 is easy and straightforward. We use the original image resolution and pad it to 1280×384 . We train the network with a batch size of 32 on 4 GeForce TITAN X GPUs for 60 epochs. The learning rate is set at 2.5×10^{-4} and drops at 25 and 40 epochs by a factor of 10. During testing, we use the top 100 detected 3D projected points and filter it with a threshold of 0.25. No data augmentation method and NMS are used in the test procedure. Our implementation platform is Pytorch 1.1, CUDA 10.0 and CUDNN 7.5.

Depth Estimation from Monocular Images and Sparse Radar Data

In this page, we explore the possibility of achieving a more accurate depth estimation by fusing monocular images and Radar points using a deep neural network. We give a comprehensive study of the fusion between RGB images and Radar measurements from different aspects and proposed a working solution based on observations. We find that the noise existing in Radar measurements is one of the main key reasons that prevents one from applying the existing fusion methods developed for LiDAR data and images to the new fusion problem between Radar data and images. The experiments are conducted 实施 on the nuScenes dataset, which is one of the first datasets which features Camera, Radar, and LiDAR recordings in diverse 各种各样的 scenes and weather conditions. Extensive 广泛的 experiments demonstrate that our method outperforms existing fusion methods. We also provide detailed ablation 灼烧 studies to show the effectiveness of each component in our method.

Dense and robust depth estimation is an important component in self-driving system and unmanned 无人 aerial vehicles. While existing structure-light-based depth sensor or stereo camera can provide dense depth in indoor environments, the reliability of these sensors degrade a lot in outdoor applications. As a result, lots of research works focus on obtaining dense depth from monocular RGB images only. Recently, convolutional neural network (CNN) based methods have demonstrated impressive improvement on monocular depth estimation for both indoor and outdoor scenarios. However, there is still a gap between the accuracy and reliability of these methods and what the real-world applications need.

Apart from estimating depth from monocular camera, to improve the robustness of the system, some methods also take other sensor modalities形式 into consideration. While these sensors, LiDAR is the most commonly used one. Many works have been conducted on dense depth estimation from RGB images and sparse LiDAR scans. In addition to depth estimation and completion tasks, different RGB + LiDAR fusion techniques are also extensively used in tasks such as 3D object detection. Although LiDAR provides more accurate depth measurements in outdoor scenario, high-end LiDAR sensors are still far from affordable for many applications.

Compared with LiDAR, Radar is an automotive-grade sensor that has been used for decades on vehicles, but has not attracted lots of attention in self-driving research based on deep learning. One reason might be that Radar measurements are not included in most of the dominant self-driving datasets. Compared with LiDAR, Radar sensors offer longer sensing range (200m~300m), more attributes including velocities, dynamic states, and measurement uncertainties. Most importantly, the costs of these sensors are much lower than LiDAR. However, Radar measurements are much lower than LiDAR, However, Radar measurement are typically sparser, noisier, and have a more limited vertical field of view.

This work is to study the challenges of using Radar data for dense depth estimation and to propose a novel method for that aim. Given recently released nuScenes dataset consisting of RGB, LiDAR and Radar measurements, we are able to conduct experiments on cross-modality sensor fusion between RGB camera and Radar. Through our experiments we demonstrated that: 1) Existing RGB + LiDAR fusion methods can be applied directly to RGB + Radar fusion task; and 2) with proper fusion strategies and a novel denoising operation, our proposed network is able to improve the performance of depth estimation by a good margin by using Radar measurements. According to our survey, our work is the first one that brings Radar sensors into dense depth estimation tasks.

The contributions of this work include: 1) a detailed study on the challenges of using Radar data for dense depth estimation; and 2) a novel and carefully motivated network architecture for depth estimation with monocular images and sparse Radar data.

RELATED WORKS

RGB-based Depth Estimation. Depth estimation from monocular or stereo camera is a popular research topic in both computer vision and robotics. Early works used either geometry-based algorithms on stereo images, or handcrafted features on single images. Recent years, convolutional neural networks (CNN) have demonstrated their ability in image understanding, dense predictions, etc. given large scale datasets. Therefore, lots of research works of monocular depth estimation are conducted. In general, most of them used the encoder-decoder architectures. Xie further introduced skip connection strategy which is a frequently used technique to multi-level features in dense prediction tasks. On the other hand, Huang achieve state-of-the-art performance by introducing space increasing discretization (SID) and ordinal regression. In some semi-/self-supervised formu loss is used.

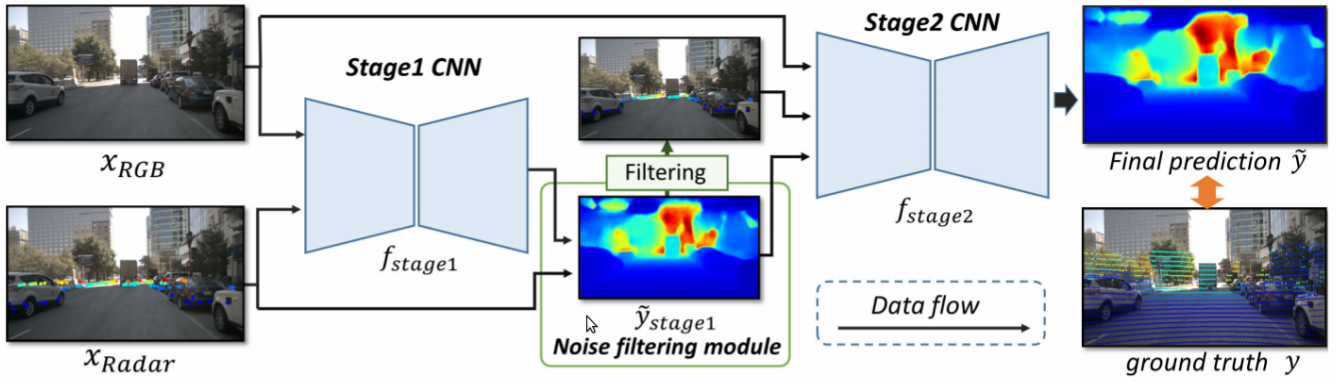


Fig. 1: The full pipeline of our method. Each CNN stage is a full late fusion model described in Section III-B.

and the smoothness constraint 约束 is further imposed to enhance local consistency 一致性. Patil proposed a recurrent 反复发作 network architecture to exploit the long-range spatiotemporal 时空 structures across video frames to yield more accurate depth maps. While good performance has been obtained with only RGB images, the methods still have difficulty in generalizing to new scenario and challenging weather and light conditions. This motivates the existing line of work that fuses camera data with Lidar data and our work that fuses camera data with radar data is cheaper to obtain.

Depth Estimation with Camera and Lidar data

while monocular depth estimation task attracts lots of attention, achieving more reliable and accurate predictions using multi-modality 多模态 information is also a popular topic. Existing works either take the whole set of LiDAR points, (known as depth completion), or the downsampled set as model inputs. Ma first projected LiDAR points to 2D sparse depth map and then perform so called early fusion by direct concatenation with RGB images along channel, or concatenation feature maps after one shallow convolution block. J used a late fusion method to combine features from different modalities and improved the overall performance through multitask learning. Q proposed to predict dense depth map by combining predictions from RGB and surface normal pathways, where surface normal is treated as an intermediate representation. Moreover, confidence maps are predicted to down-weight mixed measurements from LiDAR caused by the displacement between camera and LiDAR. In this work our main focus is sensor fusion. Thus, we use the widely adopted encoder-decoder architecture and focus on the necessary extensions in order to effectively use Radar data instead.

Post-processing and Refinement Methods 后处理和优化方法. Apart from treating sparse point clouds as inputs to the model, some methods also tried to directly refine the dense predictions of the trained models. W proposed a simple add-on module that can improve the prediction of depth estimation model using similar methods used by white box adversarial attack. Since the refinement is done using iterative re-inference, no re-training is required. This method can be integrated into most deep learning based methods. Cheng learned an affinity 亲和力 matrix from data to refine the outputs of their CNN model. The recurrent refinement operation can also be extended to depth completion tasks.

Fusion of images and radar Data.

There are already works that fuse RGB images and Radar. given the fact that they are very much complementary 补充. The line of work mainly focus on object detection and tracking. For instance, C fused 融合 Radar data and images to detect small objects at a large distance. In [] and [] the authors enhance current 2D object detection networks by fusing camera data and projected sparse radar data in the network layers, while [] also performs free space semantic segmentation jointly. Both methods learn at

which level the fusion of the sensor data is more beneficial for the task. There are also other datasets proposed for object detection with Radar data such as []. Exemplary 示范性的 works on semantic segmentation with Radar point cloud have been conducted as well. For instance, in [] the authors have studied how the challenging task can be performed and provide results on a dataset with manually labeled radar reflections. Similar to these works for object detection and semantic segmentation, our work aims to study how the challenging task of dense depth estimation with radar data can be addressed with the popular deep neural network architectures.

In the line of increasing the robustness of depth estimation, Vasudevan, et al. [] have proposed a novel method to estimate depth maps based on binaural 双耳 sounds.

Our Method

Our whole method are divided into multiple main components. In the following section, we will go through each component in detail.

Radar Data background

Different from well established depth completion or depth estimation tasks. There is no prior research works on RGB + radar depth estimation task. Therefore, we provide a brief introduction to the task formulation and some key differences between Radar and LiDAR measurements, which will help readers to understand the motivations behind the components of our method.

Data format. similar to LiDAR data, Radar measurements are recorded as sparse point clouds. The main difference is that, in addition to x, y, z , and reflectance , Radar data consist of additional measurements including the velocity along x and y direction, the standard deviation of the location and velocity measurements, and information such as the dynamic states of the measured object (encoded as discrete 离散 numbers).




Fig. 2: Limitations of Radar measurements. (Left) Noisy measurements and (Right) Limited vertical field of view.

Limitation . While it seems that the radar data provide more information, it also introduces the following limitations compared with LiDAR data:

- sparseness: In nuScenes dataset, there are more than 3000 LiDAR points after projection to the camera. However, there are less than 100 Radar points after the projection.
- Limited vertical field of view. Because of the limitation of the sensor, Radar measurements mainly concentrate in the central horizontal bin (similar heights) as shown in Fig.2.
- Noisy measurements: Due to multiple reflections (Radar multipath problem) or other reasons, we have many noisy measurements as shown in Fig.2.
- Inconsistency 前后不一致 with LiDAR data: Apart from noisy measurements, which are considered as outliers, the 3D points of Radar and LiDAR representing the same object can also be different. Since we typically use LiDAR measurements as ground truth, even noise-free Radar measurements are not perfect on the evaluation metrics.

As we will show in Section IV, using Radar depth maps directly as the input of off-the-shelf 现成的 RGB+LiDAR depth completion /prediction models resulted in marginal 边缘 improvements.

Problem formulation. In our RGB+ Radar formulation, each data sample from the dataset contains

- (1) an RGB image x_{RGB}
- (2) a set of Radar measurements $R = \{r_n\}_{n=1}^N$ from 3 nearest timestamps
- (3) a set of Lidar measurements $L = \{l_m\}_{m=1}^M$

Radar measurements R can be further projected to a single-channel 2D depth map x_{Radar} using the perspective projection. Similarly, LiDAR measurements can be objected to 2D map y , which is treated as ground truth depth map in our experiments (section IV). Our model takes both x_{RGB} and x_{Radar} as inputs and predicts dense 2D depth map \tilde{y} which minimizes the metrics errors. Same as all the depth estimation /completion tasks, loss and metric error are computed over the pixels with ground truth measurements.

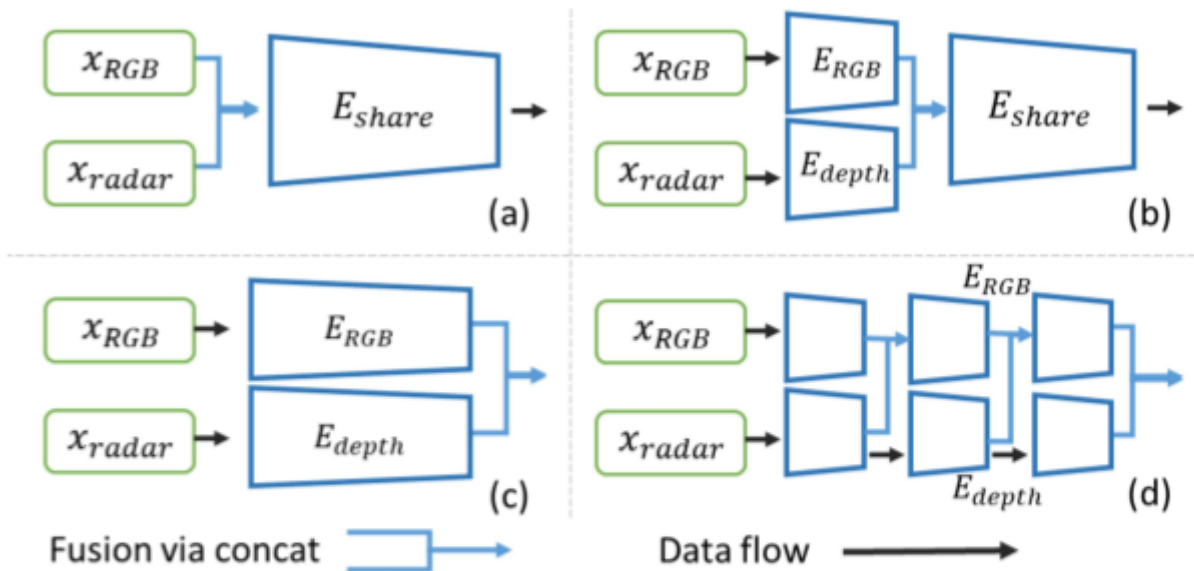


Fig. 3: Fusion methods for encoder: (a) Early fusion, (b) Mid fusion, (c) Late fusion, and (d) Multi-layer fusion. Due to space limit, we didn't show details layer-wise.

CNN architecture

Backbone. Since our goal is to perform a comprehensive综合的 study on different fusion methods, we need to choose an efficient and widely-used backbone. Thus, we fixed our backbone to Resnet18 , and explored different fusion methods based on it as a pilot飞行员 experiment. As illustrated in Fig.3, we apply different encoder fusion methods to the simple encoder-decoder architecture proposed by Ma, and compare their performance. According to the experiment (Section IV-C), late fusion and multi-layer fusion model have comparable performance. Therefore, we adopt采纳 late fusion as our main encoder backbone design in the following experiments for simplicity.

Two-Stage architecture As shown by the pilot experiments(section IV-C) , a simple encoder-decoder architecture have some improvements on RGB + Radar depth prediction task if we can remove most of the noisy measurements. However, we don't have LiDAR ground truth to help us performing the filtering in real applications.and it's hard to perform outlier rejections异常剔除 without information on the 3D structure or objects of the scene. Therefore, we come up with a 2-stage design to address the noisy measurement issue.

As shown in Fig.1, our whole method contains two stages. The stage 1 model f_{stage1} takes both RGB image x_{RGB} and the radar depth map x_{Radar} as inputs and predicts a coarse粗 depth map $\tilde{y}_{\text{stage1}}$, which gives us a dense 3D structure of the scene:

$$\tilde{y}_{\text{stage1}} = f_{\text{stage1}}(x_{\text{RGB}}, x_{\text{Radar}}) \quad (1)$$

Then we compare the Radar depth map with the coarse prediction $\tilde{y}_{\text{stage1}}$ to reject some outliers (most details in next subsection) and obtain the filtered Radar depth map \tilde{x}_{Radar} . The assumption假设 here is that although the predictions from stage 1 is not perfect, they are smooth and locally consistent. Therefore, they are suitable to reject outlier noises produced by Radar Multipath, which typically have certain margins with the correct depth values.

The stage2 model f_{stage2} takes x_{RGB} , \tilde{x}_{Radar} , and the prediction from stage1 $\tilde{y}_{\text{stage1}}$ to predict the final result \tilde{y} :

$$\tilde{y} = f_{\text{stage2}}(x_{\text{RGB}}, \tilde{x}_{\text{Radar}}, \tilde{y}_{\text{stage1}}) \quad (2)$$

Noise filtering module

Since Radar measurements are not exactly consistent with the LiDAR measurements as we mentioned in Section III-A .we need to have some tolerance公差 in the filtering process. Otherwise , we will end up discarding丢弃 all the measurements in the set R .

Instead of setting a fixed distance tolerance threshold τ , we empirically凭经验 found that an adaptive threshold gives us better results. We design the threshold to be a function of depth value $\tau(d)$: We have larger tolerance for larger depth values, which is similar to the space-increasing discretization (SID) from Huan et al.[2]:

$$\tau(d) = \exp\left(\frac{d \cdot \log(\frac{\beta}{\alpha})}{K} + \log(\alpha)\right) \quad (3)$$

here we heuristically启发式的 set $\alpha=5$ and $\beta=18$.

Let P denote the set pixel coordinates (u,v) of the Radar measurements projected by the perspective projection function $\text{proj}(\cdot)$: $P = \{p_n \in N_{n=1} = \{\text{proj}(r_n) \in N_{n=1}\}$. The noise filtering module will keep the point p_n if it satisfies the following constraint约束:

$$\|x_{\text{Radar}}(p_n) - \tilde{y}_{\text{stage1}}(p_n)\| \leq \tau(p_n), \text{ for } p_n \in P \quad (4)$$

Loss Function

By design, each component of our model is differentiable. Thus, our whole model is end-to-end trainable. Following the setting from [7], We apply L_1 loss to both the predictions of stage1 ($\tilde{y}_{\text{stage1}}$) and stage2 (\tilde{y}). Considering that the main purpose of $\tilde{y}_{\text{stage1}}$ is to filter outlier noises in x_{Radar} , we further add edge-aware smoothness constraint to it. To effectively balance multiple loss terms, we follow the method proposed by Kendall

$$L_{\text{total}} = e^{-w_1} (L_1(\tilde{y}_{\text{stage1}}, y) + 10^{-3} * L_{\text{smooth}}) + e^{-w_2} * L_1(\tilde{y}, y) + \sum_i w_i \quad (6)$$

where w_1 and w_2 are optimized variables, and L_{smooth} is defined as:

$$L_{\text{smooth}} = \frac{\nabla_u(\tilde{y}_{\text{stage1}})}{e^{-\nabla_u(x_{\text{RGB}})}} + \frac{\nabla_v(\tilde{y}_{\text{stage1}})}{e^{-\nabla_v(x_{\text{RGB}})}} \quad (8)$$

∇_u and ∇_v denote the gradient along 2D height and width directions separately.

Implementation details

Unless stated otherwise, all the models are trained using a batch size of 16 and SGD optimizer with a learning rate of 0.001 and a momentum 势头 of 0.9 for 20 epochs. The learning rate is multiplied by 0.1 after every 5 epochs. All the models we used in experiment action are implemented in PyTorch. The experiments are conducted on the desktop computers/clusters with Nvidia GTX1080Ti and TeslaV100 GPUs.

All of our encoder network architectures (section III-B) are modified from the standard Resnet18. For early fusion, we simply modified the input channels to 4 and randomly initialized the weights (weights of other layers were initialized from pre-trained models on Imagenet dataset). For mid, late, and multi-layer fusion, the depth branch has a similar architecture as the RGB branch. The only difference is that we change the number of channels to 1/4 of the original one. Fusion operations happened only on feature maps with same spatial resolution (width and height). Regarding the decoder part, we kept the section from [] by using UpProj module as our upsampling operation.

Experiment